

Theo Song
Prof. Hakner
ECE 357
10/30/2019

Unix: CentOS Architecture: x86_64 Kernel: 3.10

Problem 1 - Signal Numbers & Behavior

1) SIGTSTP

2) Ctrl + \

3) stty intr ^I

```
4) sigset_t set;
   if(sigemptyset(&set)<0)
       perror("sigemptyset\n");
   if(sigaddset(&set, 60)<0)
       perror("sigaddset\n");
   if(sigprocmask(SIG_UNBLOCK, &set, NULL)<0)
       perror("sigprocmask\n");
```

For signal #60, the handler function will run 32 times. Because signal #60 is a real time signal, meaning that it could be queued, the handler function will run as much as the signal is sent.

However, if it was signal#2, the handler function will run only once. Since signal #2 is a traditional signal, when a new signal(#2) is sent, it will overwrite the older signal, but will not queue. Thus, in the case of signal #2, the handler function will run only once.

Problem 2 - File Descriptor Tables

- 1) It is because some of the `write()` syscall are actually successful, meaning that it fully wrote 65536 bytes into the pipe buffer, which will not print the pipe short write if such is the case. In order for the `write()` to be successful, a `signal(SIGUSR1)` must not change the `write()` into a `READY` state, while the `write()` is in the interruptible `SLEEP` state. The pipe short write is written when the `SIGUSR1` interrupts the `write()` in `SLEEP` state, which has written some bytes of 4K chunks. When the `write()` is changed to the `READY` state due to the signal, not actually being ready, it returns the bytes it has written so far, which is the number shown in “pipe short write”.
- 2) All the short write values are a multiple of 4096(4K). This is because when the requested write size exceeds 4K, the write is separated into smaller chunks of 4K, which could be interrupted by other writes or signal.
- 3) `EINTR` is shown in some parts of the `stderr`. While the write is in the `SLEEPING` state, waiting for the `read()` to read the pipe buffer, the `SIGUSR1` signal interrupts, transitioning `write()` to the `READY` state. Once scheduled, the kernel notice that it has been interrupted by a signal. Now, if the `write()` has done no work, and the `SA_RESTART` flag is set, the `write()` syscall is restarted after the `f()` `sig_handler` returns. However, without the `SA_RESTART` flag, the `write()` will be considered a failure and `errno` will be set to `EINTR`, after the `f()` `sig_handler` returns. And if the `write()` has done partial work, it will return the partial work it has done. That is why some shows `EINTR`, while other lines show the short write value.
- 4) The program will not terminate. When a child process terminates, a `SIGCHLD` is sent to the parent. The `signal(SIGCHLD, SIG_IGN)` is saying that the parent has no interest of the child, thus not making a zombie. However, without such statement, when the child process (w & r) terminates, it will remain in the `ZOMBIE` state, waiting for the parent process to claim the status, using the `wait` family of system calls. And because the parent process does not have any such, the child process (w) will remain in the `ZOMBIE` state, allowing the parent process to run the infinite loop of sending `SIGUSR1`, without an error, to the child process(w).